

CS 215

Data Analysis and Interpretation

Multivariate Statistics: Multivariate Gaussian

Suyash P. Awate

Multivariate Gaussian – Definition

- Consider a vector random variable $X := [X_1; X_2; \dots; X_D]$
 - Column vector of length D

Definition: The RV X has a multivariate (jointly) Gaussian PDF if \exists a finite set of i.i.d. univariate standard-normal RVs W_1, \dots, W_N (with $D \leq N$) such that each X_d can be expressed as $X_d = \mu_d + \sum_n A_{dn} W_n$ (i.e., $X = AW + \mu$).

Multivariate Gaussian – Identity A

- Consider a vector random variable $X := [X_1; X_2; \dots; X_D]$
 - Column vector of length D

Definition: The RV X has a multivariate (jointly) Gaussian PDF if \exists a finite set of i.i.d. univariate standard-normal RVs W_1, \dots, W_N (with $D \leq N$) such that each X_d can be expressed as $X_d = \mu_d + \sum_n A_{dn} W_n$ (i.e., $X = AW + \mu$).

- Example 1 (Zero-Mean + Isotropic / Spherical Gaussian): The case of independent standard-normal RVs W_1, \dots, W_D with $A := I_{D \times D}$ and $\mu := 0$, i.e. $X = W$

Then, the Gaussian PDF is $p(w) = \prod_d p(w_d) = \frac{1}{(2\pi)^{D/2}} \exp(-0.5w^\top w)$

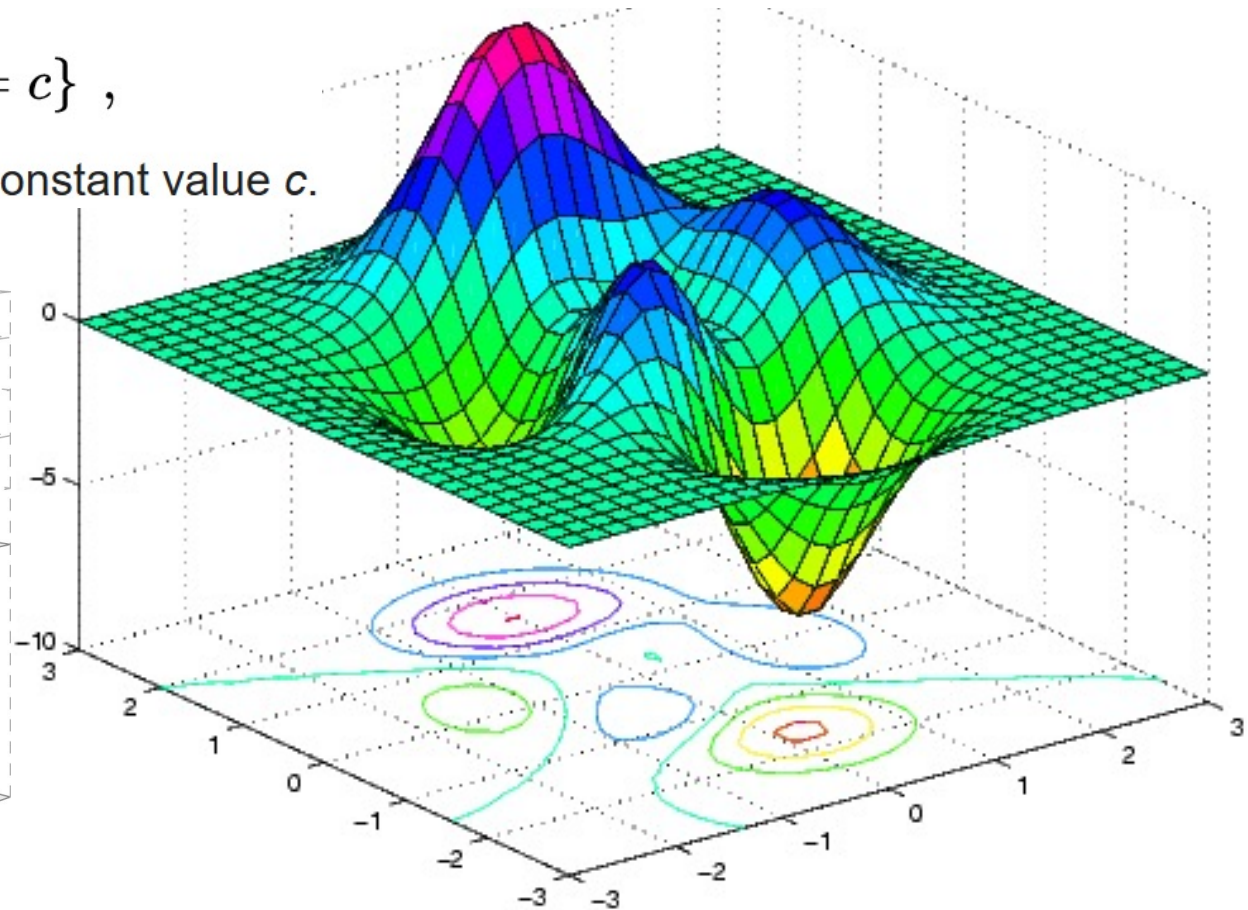

Multivariate Gaussian – Identity A

- What are the level sets of the PDF ?

In **mathematics**, a **level set** of a **real-valued function** f of n **real variables** is a set of the form



$$L_c(f) = \{(x_1, \dots, x_n) \mid f(x_1, \dots, x_n) = c\} ,$$

that is, a set where the function takes on a given constant value c .





Multivariate Gaussian – Identity A

- Isotropic / spherical multivariate Gaussian
 - Level sets




$$p(w) = \prod_d p(w_d) = \frac{1}{(2\pi)^{D/2}} \exp(-0.5w^\top w)$$



Multivariate Gaussian – Identity A

- Isotropic / spherical multivariate Gaussian
 - Level sets

$$p(w) = \prod_d p(w_d) = \frac{1}{(2\pi)^{D/2}} \exp(-0.5 w^\top w)$$






Multivariate Gaussian – Diagonal A

- $X = A W + \mu$
- What is PDF $q(X)$ for **non-singular** square **diagonal** matrix A , some μ ?
 - $X_1 = A_{11} W_1 + \mu_1$: Gaussian with mean μ_1 , standard deviation $\sigma_1 = |A_{11}|$
 - $X_2 = A_{22} W_2 + \mu_2$: Gaussian with mean μ_2 , standard deviation $\sigma_2 = |A_{22}|$
 - ...
 - $X_D = A_{DD} W_D + \mu_D$: Gaussian with mean μ_D , standard deviation $\sigma_D = |A_{DD}|$
 - $P(X) = P(X_1, X_2, \dots, X_D) = G(X_1; \mu_1, \sigma_1^2) G(X_2; \mu_2, \sigma_2^2) \dots G(X_D; \mu_D, \sigma_D^2)$
 - Any level set of PDF $q(X)$ is a hyper-ellipsoid with:
 - Center at μ
 - Axes aligned with cardinal axes

Multivariate Gaussian – Diagonal A

- $X = A W + \mu$
- What is PDF $q(X)$ for **non-singular** square **diagonal** matrix A , some μ ?
 - $P(X) = P(X_1, X_2, \dots, X_D) = G(X_1; \mu_1, \sigma_1^2) G(X_2; \mu_2, \sigma_2^2) \dots G(X_D; \mu_D, \sigma_D^2)$
 - Example 1-3 (left to right):
both means (μ_1, μ_2) are zero,
both variances are (σ_1^2, σ_2^2) : (4,4), (9,1),(1,9)




Multivariate Gaussian – Non-Singular A



- $X = A W + \mu$
- What is PDF $q(X)$ for **non-singular square** matrix A and $\mu = 0$?
- Transformation of random variables (multivariate case)
 - Transformation is $X := g(W) := A W$
 - Inverse transformation is $W = g^{-1}(X) = A^{-1}X$
 - Univariate case
 - We wanted magnitude of derivative of $g^{-1}(.)$
 - Measured local scaling in lengths caused by $g^{-1}(.)$
 - Multivariate case
 - Measure local scaling in volumes caused by $g^{-1}(.)$
 - We want the magnitude of the volume-scaling given by Jacobian of $g^{-1}(.)$
 - Magnitude of determinant of Jacobian of $g^{-1}(.)$

Multivariate Gaussian – Non-Singular A

- Linear transformation
 $W := A^{-1} X$




T_M



Multivariate Gaussian – Non-Singular A

- Linear transformation $W := A^{-1} X$
 - Transformation A^{-1} maps
an infinitesimal hyper-cube (dX) $\delta \times \delta \times \dots \times \delta$ (D times) \rightarrow
an infinitesimal hyper-parallelepiped (dW)
 - If axes of hyper-cube were unit vectors along cardinal axes,
then axes of hyper-parallelepiped are columns of A^{-1}
 - If volume of the hyper-cube (dX) is δ^D ,
then volume of hyper-parallelepiped (dW) is $\delta^D \det(A^{-1}) = \delta^D / \det(A)$




Multivariate Gaussian – Non-Singular A

- Volume of a parallelepiped (in 3D)
 - Scalar triple product

$$\begin{aligned}\mathbf{a} \cdot (\mathbf{b} \times \mathbf{c}) &= \begin{vmatrix} a_x & a_y & a_z \\ b_x & b_y & b_z \\ c_x & c_y & c_z \end{vmatrix} = \begin{vmatrix} a_x & b_x & c_x \\ a_y & b_y & c_y \\ a_z & b_z & c_z \end{vmatrix} \\ &= \mathbf{b} \cdot (\mathbf{c} \times \mathbf{a}) = \mathbf{c} \cdot (\mathbf{a} \times \mathbf{b}) \\ &= -\mathbf{a} \cdot (\mathbf{c} \times \mathbf{b}) = -\mathbf{c} \cdot (\mathbf{b} \times \mathbf{a}) = -\mathbf{b} \cdot (\mathbf{a} \times \mathbf{c})\end{aligned}$$

The notation $[\mathbf{a}, \mathbf{b}, \mathbf{c}]$ is also used for $\mathbf{a} \cdot (\mathbf{b} \times \mathbf{c})$.






$$\begin{aligned}\text{Volume} &= \text{area of base} \cdot \text{height} \\ &= \|\mathbf{a} \times \mathbf{b}\| \|\mathbf{c}\| |\cos \phi| = |(\mathbf{a} \times \mathbf{b}) \cdot \mathbf{c}|\end{aligned}$$

Multivariate Gaussian – Non-Singular A

- Why is volume of hyper-parallelepiped given by determinant of matrix with columns as sides of hyper-parallelepiped ?
 - The following is an argument (not a proof; a separate inductive proof exists):
 - 2 important properties from linear algebra:
 - Adding multiples of one column/side vector to another:
 1. doesn't change determinant, because determinant function is multi-linear
 2. doesn't change volume, because it causes a skew translation of hyper-parallelepiped
- Using Gram-Schmidt orthogonalization, transform matrix A^{-1} to a matrix, say, A^{-1}_{ortho} with orthogonal columns (NOT orthonormal columns; that would have determinant 1)
 - This doesn't change determinant or volume

Multivariate Gaussian – Non-Singular A

- Gram–Schmidt orthogonalization
 - $\{v_1, v_2\}$ to $\{u_1, u_2\}$







Multivariate Gaussian – Non-Singular A

- Why is volume of hyper-parallelepiped given by determinant of matrix with columns as sides of hyper-parallelepiped ?
 - The following is an argument (not a proof; a separate inductive proof exists):
 - 2 important properties from linear algebra:
Adding multiples of one column/side to another:
 - 1) doesn't change determinant, because determinant function is multi-linear
 - 2) doesn't change volume, because it causes a skew translation of hyper-parallelepiped
- Using Gram-Schmidt orthogonalization,
transform matrix A^{-1} to a matrix, say, A^{-1}_{ortho} with orthogonal columns
(NOT orthonormal columns; that would have determinant 1)
- Rotate A^{-1}_{ortho} to make it to diagonal form (align columns to cardinal axes)
 - This doesn't change determinant or volume

Multivariate Gaussian – Non-Singular A

- Rotation / alignment to cardinal axes



Multivariate Gaussian – Non-Singular A

- Why is volume of hyper-parallelepiped given by determinant of matrix with columns as sides of hyper-parallelepiped ?
 - An intuitive argument (not a proof; a separate inductive proof exists):
 - Adding multiples of one column/side to another:
 - 1) doesn't change determinant, because determinant function is multi-linear
 - 2) doesn't change volume, because it causes a skew translation of hyper-parallelepiped
 - Using Gram-Schmidt orthogonalization, transform matrix A^{-1} to a matrix, say, A^{-1}_{ortho} with orthogonal columns (NOT orthonormal columns; that would have determinant 1)
 - Rotate A^{-1}_{ortho} to make it to diagonal form (align columns to cardinal axes)
 - For this diagonal matrix (aligned hyper-rectangle), determinant magnitude (= product of diagonal-entries' magnitudes) = volume of a hyper-rectangle (= product of side lengths)
 - Now trace back all operations

Multivariate Gaussian – Non-Singular A

- $X = A W + \mu$
- What is the PDF $q(X)$ for non-singular square matrix A and $\mu = 0$?
- Transformation of random variables (multivariate case)
 - Transformation is $X := g(W) := A W$
 - Inverse transformation is $W = g^{-1}(X) = A^{-1}X$
 - Multivariate case
 - Measure local scaling in volumes caused by $g^{-1}(.)$
 - We want the magnitude determinant of Jacobian of $g^{-1}(.)$

$$q(X) = p(A^{-1}X) \frac{1}{|\det(A)|} = \frac{1}{(2\pi)^{D/2} |\det(A)|} \exp(-0.5 X^\top (A^{-1})^\top A^{-1} X)$$

Let $C := A A^\top$. Then, $C^{-1} = (A^{-1})^\top A^{-1}$ and $\det(C) = \det(A) \det(A^\top) = (\det(A))^2$

$$q(X) = \frac{1}{(2\pi)^{D/2} |C|^{0.5}} \exp(-0.5 X^\top C^{-1} X)$$

Multivariate Gaussian – Non-Singular A, Non-Zero μ

- If $X = A W$ is a multivariate Gaussian,
then $Y = X + \mu$ is a multivariate Gaussian with

$$p(y) = \frac{1}{(2\pi)^{D/2} |C|^{0.5}} \exp(-0.5(y-\mu)^\top C^{-1}(y-\mu))$$

- Proof:
 - Follows from the transformation $X := Y - \mu := g^{-1}(Y)$

Multivariate Gaussian – Composite Transformations

- If Y is multivariate Gaussian, then $Z := BY + c$ is multivariate Gaussian, where matrix B is square invertible
- Proof:
 - Because Y is multivariate Gaussian, we have $Y = AW + \mu$, where A is invertible
 - Thus,
$$\begin{aligned} Z &= B(AW + \mu) + c \\ &= (BA)W + (B\mu + c), \text{ where matrix } BA \text{ is invertible} \end{aligned}$$

Multivariate Statistics – Mean and Covariance

Multivariate Statistics – Mean

- For an general random (column) vector X , the mean vector is

$$E_{P(X)}[X]$$

= a (column) vector with the i -th component as $E_{P(X)}[X_i] = E_{P(X_i)}[X_i]$

Multivariate Statistics – Covariance

- Covariance matrix for a general random (column) vector Y :

$$C := E_{P(Y)} [(Y - E[Y]) (Y - E[Y])^T]$$

- So,

$$\begin{aligned} C_{ij} &= E_{P(Y)} [(Y_i - E[Y_i]) (Y_j - E[Y_j])] \\ &= E_{P(Y_i, Y_j)} [(Y_i - E[Y_i]) (Y_j - E[Y_j])] \\ &= \text{Cov} (Y_i, Y_j) \end{aligned}$$

Multivariate Statistics – Covariance

- More properties of covariance matrix C (for a general random vector X)

(1) $C = E[XX^\top] - E[X](E[X])^\top$

Proof: Expand the terms in the definition

(2) C is symmetric

Proof: $C_{ij} = \text{Cov}(X_i, X_j) = \text{Cov}(X_j, X_i) = C_{ji}$

(3) C is positive semi-definite (PSD)

Proof: For any $D \times 1$ non-zero vector a , we get $a^\top C a = E[a^\top (X - E[X])(X - E[X])^\top a] = E[(f(X))^\top f(X)] \geq 0$ that is the variance of a scalar RV $f(X) = (X - E[X])^\top a$

Multivariate Gaussian – Mean and Covariance

Multivariate Gaussian – Mean

- The **mean** vector of $X := AW + \mu$ is μ

- Proof:

- When $X = AW + \mu$,

$$E_{P(X)}[X] = E_{P(W)}[AW + \mu] = \mu + E_{P(W)}[AW] = \mu + A E_{P(W)}[W] = \mu$$

- Notes:

- Take the expectation of first component of AW , i.e.,

$$\begin{aligned} & E_{P(W)} [A_{11}W_1 + A_{12}W_2 + \dots A_{1D}W_D] \\ &= A_{11} E_{P(W)} [W_1] + A_{12} E_{P(W)} [W_2] + \dots + A_{1D} E_{P(W)} [W_D] \end{aligned}$$

- So, for the whole vector: $E_{P(W)} [AW] = A E_{P(W)} [W]$

Multivariate Gaussian – Covariance

- The **covariance** matrix of $X := AW + \mu$ is AA^\top

$\text{Cov}(W) = E[WW^\top] = I$ because:

(i) $\text{Cov}(W_i, W_i) = 1$ and

(ii) $\text{Cov}(W_i, W_{j \neq i}) = 0$ because of independence of W_i and W_j

$$\text{Cov}(X) = E[(X - E[X])(X - E[X])^\top] = E[(AW)(AW)^\top] = E[AWW^\top A^\top] = AE[WW^\top]A^\top = AA^\top$$

Thus, the RV $X = AW + \mu$ has covariance $C = AA^\top$, where $C_{ij} = \text{Cov}(X_i, X_j)$.